

# The `lazy.frame` Package

Bryan W. Lewis  
blewis@illposed.net

October 11, 2011

## 1 Preface

I’ve been working with some large-ish text files of comma separated values (CSV) recently. The files are each about three gigabytes with about 20 million rows. My computer has plenty of memory for R to load each file.

But, it takes a while.

And I’m impatient.

Now, I don’t really need the entire data set in memory. I really just need to filter the data a bit and then sample from the rows. I think that this situation is typical enough—wanting fast access to subsets of large text files—that I wrote this package for it.

The `lazy.frame` package lets me quickly and efficiently work with subsets from a text file without loading the entire file into memory. A “`lazy.frame`” presents a text file as a kind of simple data frame, but without first loading the file into memory. Lazy frames lazily load data from their backing files only when required, for example by an indexing operation. They are essentially lazy wrappers for the `read.table` function with a few extra convenience functions.

There are several compelling R packages for working directly with file-backed data: The [bigmemory](#) package by Emerson and Kane provides a memory mapped matrix object, free from R indexing constraints, and a comprehensive suite of fast analysis functions. The nicely simple but powerful [mmap](#) package by Jeff Ryan defines a data frame-like memory mapped object. And the venerable [ff](#) package by Adler, Oehlschlägel, et. al. defines a variety of memory mapped data frame-like objects and functions. All of these packages have really interesting features. Most of them are designed to facilitate working with objects larger than the physical RAM available on a computer.

But, recall that my data sets fit into the RAM on my computer (RAM is really cheap)! My main irritation is the bottleneck incurred by parsing the entire data set, which isn’t really avoided by the above packages (although some of the packages do include methods to help expedite loading data

from text files).

Of course, lazy frames aren't a panacea and have limitations discussed below. For *really* large data sets, or for more sophisticated operations involving all the data, `bigmemory` is a better option. Lazy frames are really good for quickly extracting subsets from large text files with between roughly a million and a hundred million or so rows.

## 2 Using `lazy.frame` package

## 3 Limitations

## 4 Examples

I present a few examples that compare indexing operations on lazy frames with indexing operations on data frames read in by `read.table`. All experiments were conducted on a 2 GHz, four-core AMD Opetron computer with 12 GB of DDR-2 RAM running Ubuntu 9.10 GNU/Linux and R version 2.12.1. The data files resided on a Fusio-io ioXtreme solid state disk rated at 700MB/s data read rate and 80 $\mu$ s read latency. In order to minimize disk caching effects between tests, the command

```
echo 3 > /proc/sys/vm/drop_caches
```

was issued (wiping clean the Linux disk memory cache) just before each test.

### 4.1 An uncompressed file example

`read.table` results:

load time:

	user	system	elapsed				
	648.380	33.350	682.699				
		used	(Mb)	gc trigger	(Mb)	max used	(Mb)
Ncells		138089	7.4	667722	35.7	380666	20.4
Vcells		285413776	2177.6	832606162	6352.3	1034548528	7893.0
[1]	17826159		27				

subset time:

	user	system	elapsed
	27.87	2.41	30.31
[1]	95166		27

`read.table` with `colClasses` results:

load time:

	user	system	elapsed				
		used	(Mb)	gc trigger	(Mb)	max used	(Mb)
Ncells	138519	7.4		350000	18.7	350000	18.7
Vcells	285348278	2177.1		649037152	4951.8	641872298	4897.1
[1]	17826159	27					

subset time:

	user	system	elapsed
	28.410	2.180	30.593
[1]	95166	27	

file.frame results:

load time:

	user	system	elapsed				
		used	(Mb)	gc trigger	(Mb)	max used	(Mb)
Ncells	140517	7.6		350000	18.7	350000	18.7
Vcells	130910	1.0		786432	6.0	531925	4.1
[1]	17826159	27					

subset time:

	user	system	elapsed
	40.870	11.770	52.709
[1]	95166	27	